



Emotion and Affect Representation in Sentence Embeddings

Author:

Luis Alberto BARRADAS
CHACÓN
278183

Supervisors:

Prof. Dr. Christian
WARTENA
MSc. Rafael DRUMOND

2nd May 2020

Thesis submitted for
MASTER OF SCIENCE IN DATA ANALYTICS

WIRTSCHAFTSINFORMATIK UND MASCHINELLES LERNEN
STIFTUNG UNIVERSITÄT HILDESHEIM
UNIVERSITÄTSPLATZ 1, 31141 HILDESHEIM

Statement as to the sole authorship of the thesis:

Emotion and Affect Representation in Sentence Embeddings.

I hereby certify that the master's thesis named above was solely written by me and that no assistance was used other than that cited. The passages in this thesis that were taken verbatim or with the same sense as that of other works have been identified in each individual case by the citation of the source or the origin, including the secondary sources used. This also applies for drawings, sketches, illustration as well as internet sources and other collections of electronic texts or data, etc. The submitted thesis has not been previously used for the fulfilment of a degree requirements and has not been published in English or any other language. I am aware of the fact that false declarations will be treated as fraud.

2nd May 2020, Hildesheim

Abstract

Word embeddings exist to abstract features of text into a numeric space. In this project we explore the representation of different emotions and affects from established labeled datasets in common word embeddings.

Contents

1	Acknowledgements	1
2	Introduction	2
2.1	Emotions and Affect	2
	Models of Emotions	2
2.2	Emotion and Machine Learning	3
2.3	Problem setting	3
2.4	Project Description	3
2.5	Objective	3
2.6	Justification	3
2.7	A Section	4
	A subsection	4
3	Related Work	6
4	Methodology	9
5	Experiments	11
6	Conclusion	12

List of Figures

2.1	Example of an image	5
2.2	Example of an equation	5

List of Tables

Listings

Chapter 1

Acknowledgements

Chapter 2

Introduction

2.1 Emotions and Affect

Within the context of this project it is important to distinguish between emotion and affect. Affect, in the context of this project will be treated as a term to associate predisposition towards stimuli. Thus, affect is in a sense, a general term that can be even used to describe animal, and other non-human entities. Emotions, on the other hand, are treated in this project as a state inherent to humans. This state is multidimensional, and every dimension, or emotion, can either be present in a certain amount, or not be present at all. Emotions present an affect value, but not necessarily otherwise.

Models of Emotions

When trying to detect emotion, it is relevant to know what emotions to look for. This is called a model of emotions, and is still a very discussed subject in psychology. Although there are several models of emotions the most persistent are Ekman's model, Plutchik's model. The main assumption this work follows is the theory of constructed emotions, which recognizes affect as a physiological response to positive or negative stimuli, but emotions as a cognitive form of context-giving.

Why is it important to study emotion? Emotions are considered a human state that influences behaviour and decision making. Many times, when expressing thoughts in a written or spoken form, one or several emotions are present. Detecting these emotions is an important task for human interaction. Automatic emotion detection on text is thus a machine learning task required for comprehensive human-computer interaction.

Eckman's model of Emotions

What are emotions? What is affect? Why study emotions? The role of emotions in communication Measurements of emotions: - Facial Expressions - Biosignals - Language or self report Constructed theory of emotions Emotions in Language Semantic Fields Emotion Lexicon Emotion Networks Learning word semantics from context

2.2 Emotion and Machine Learning

Language in Machine Learning NLP workflow? Language representation Tokenization The problem of large dictionaries Dimensionality reduction/autoencoders Deep learning for language: automatic feature extraction Machine Learning meaning from context Word Embeddings Transformers

2.3 Problem setting

2.4 Project Description

2.5 Objective

2.6 Justification

Word Embeddings Numerically representing words is commonly known as word embedding. This allows for Machine Learning (ML) models to easily manipulate text data that would otherwise be an arbitrary encoding of text. With the advances in machine learning automatic word embedding became a possible solution to avoid crowdsourcing. Several machine learning approaches try to automatically learn the best numeric representation for characters, words, or sentences given the context of a dataset. Recently there have been many efforts from research institutions to generalize these embeddings through the use of powerful models, and bigger datasets. Such is the case of BERT, a model created by Google with massive datasets.

Affect learning and its relevance The relevance of affect in text has increased since the popularization of text-based social networks, like twitter. There, individuals and organizations openly express their opinions. This creates an environment where implicit feedback about entities is present. An easy way to abstract popular opinion about a named entity is learning the affect expressed in text, such as a tweet. Affect can be a multidimensional

phenomenon, but the most important dimension of it is valence: whether a text expresses positive or negative emotion.

In Figure 2.1 you can see an example of an image.

<http://saifmohammad.com/WebPages/NRC-Emotion-Lexicon.htm>

2.7 A Section

A subsection

A sub-subsection



Figure 2.1: Example of an image

$$\begin{aligned}
 A = \{ & 'a' : ['d', 'e', 'f', 'g'], \\
 & 'b' : ['d', 'e', 'f', 'g'], \\
 & 'c' : ['d', 'e', 'f', 'g'], \\
 & 'd' : ['h'], \\
 & 'e' : ['h'], \\
 & 'f' : ['h'], \\
 & 'g' : ['h'], \\
 & 'h' : [] \}
 \end{aligned}
 \tag{2.1}$$

Figure 2.2: Example of an equation

Chapter 3

Related Work

Mohamad and Turney created an Emotion Lexicon through crowdsourcing [Mohammad and Turney, 2013]. In this way an emotional word embedding was created by subjectively asking participants whether or not a word was related to a specific emotion.

Vo and Zhang created an automatic approach to learning sentiment lexicons for short texts through the use of emojis [Vo and Zhang, 2016]. This method uses the intrinsic usage of emojis to express positive or negative valence in a sentence, and exploded this to expand that valence to words used in the same context.

Maas et al. created a method to learn word vectors for sentiment analysis [Maas et al., 2011].

By applying machine learned automatic embeddings, the creation of word embeddings based only on text data was open as a possibility. This is also a method that later became the popular Word2Vec method [Mikolov et al., 2013].

A refining of word embeddings has been suggested by Yu et al. by means of a clustering algorithm on the vector space [Yu et al., 2017].

Rothe et al. suggested an orthogonal transformation to word embeddings used on SemEval2015 which yielded ultradense word embeddings for affect [Rothe et al., 2016].

A further exploration of transformations of a word vector space was done by Hollis et al. by means of component analysis, thus creating models of semantics from text. These were applied to affect [Hollis and Westbury, 2016].

These studies have mostly been done with affect: positive and negative valences, but have mostly ignored other emotional dimensions.

Talk about the models used and their history

Datasets There is the possibility of using the datasets from the unified dataset of emotion in text. (An Analysis of Annotated Corpora for Emotion Classification in Text). This includes:

- AffectiveText

- Blogs
- CrowdFlower
- DailyDialogs
- Electoral-Tweets
- EmoBank
- EmoInt
- Emotion-Stimulus
- fb-valence-arousal
- Grounded-Emotions
- ISEAR
- Tales
- SSEC
- TEC

The link to these datasets can be found under the github repository for the unified emotion datasets. <https://github.com/sarnthil/unify-emotion-datasets/tree/master/datasets>

During the exploratory phase of this project, a subset of these datasets will be selected, giving priority to those with a more standardized, cleaner, or easier to access data. Models For word embeddings, different pre-trained models will be used:

- Fasttext
- Word2Vec
- GloVe
- BERT

A word embedding can also be trained from data. This is considered an optional model to analyze.

The following techniques have been used:

- TSNE

- PCA

Research Question

As a general research question we propose to answer the following: When using pre-trained models for word and sentence embedding, **is the information about the emotional and affective content or context of the word or sentence represented in the vector space?**

This question can be approached in three different ways:

- Is there a direct correlation between any of the dimensions of the vector space and human-labeled emotions and affect?
- Is there a linear transformation that will yield a direct correlation to the same human-labeled emotions?
- Is there a hierarchical structure that accurately represents the embedding of said labels?

Chapter 4

Methodology

To analyze the representation of emotions in different word embeddings, a high emphasis on dimensionality-reducing visualizations is to be done. The steps to do so are the following: First, the selected datasets must be embedded in their vector space. This will allow for a faster processing of the data, but requires making the technical decision of how the word vectors will be turned into sentence vectors to share the label in the case of datasets that contain a label for every sentence. A supervised clustering must yield accuracy similar to that of a classification task. This means testing the accuracy of a classifier on the emotion recognition task. Even though this step could be seen as optional, failing to reproduce the baseline accuracy scores for this simple task might mean that the embeddings are not even capturing the basic information about the task.

A correlational analysis will be done between every dimension of the vector space and the emotions present in the dataset. This will tell us about any linear representation of emotions in the vector space. A second study must show if a linear transformation of the vector space dimensions will yield a better correlation with the emotions of the datasets. This means either LDA, or PCA.

A hierarchical clustering will be used to analyze any possible structure in the embedded dataset in relation to emotions. A final approach to this problem can be done through the study of oppositeness. These mentioned methodology is based on answering the research question with progressive approximations. It is highly unlikely that a simple embedding model represents emotions in a single dimension in a linear manner, but it is increasingly more likely that some correlation is found with a linear transformation of the aforementioned. In case these two approaches present no information about emotions, a hierarchical clustering can extract the intrinsic information of affect in emotions. Since previous works have already shown that affect can be

represented in vector spaces of word embeddings, it would be contradictory to not find a hierarchical structure of emotions in this last step.

Chapter 5

Experiments

Chapter 6

Conclusion

Bibliography

- [Hollis and Westbury, 2016] Hollis, G. and Westbury, C. (2016). The principals of meaning: Extracting semantic dimensions from co-occurrence models of semantics. *Psychonomic Bulletin & Review*, 23(6):1744–1756.
- [Maas et al., 2011] Maas, A. L., Daly, R. E., Pham, P. T., Huang, D., Ng, A. Y., and Potts, C. (2011). Learning Word Vectors for Sentiment Analysis. page 9.
- [Mikolov et al., 2013] Mikolov, T., Sutskever, I., Chen, K., Corrado, G. S., and Dean, J. (2013). Distributed Representations of Words and Phrases and their Compositionality. page 9.
- [Mohammad and Turney, 2013] Mohammad, S. M. and Turney, P. D. (2013). Crowdsourcing a Word-Emotion Association Lexicon. *arXiv:1308.6297 [cs]*. arXiv: 1308.6297.
- [Rothe et al., 2016] Rothe, S., Ebert, S., and Schütze, H. (2016). Ultradense Word Embeddings by Orthogonal Transformation. *arXiv:1602.07572 [cs]*. arXiv: 1602.07572.
- [Vo and Zhang, 2016] Vo, D. T. and Zhang, Y. (2016). Don’t Count, Predict! An Automatic Approach to Learning Sentiment Lexicons for Short Text. In *Proceedings of the 54th Annual Meeting of the Association for Computational Linguistics (Volume 2: Short Papers)*, pages 219–224, Berlin, Germany. Association for Computational Linguistics.
- [Yu et al., 2017] Yu, L.-C., Wang, J., Lai, K. R., and Zhang, X. (2017). Refining Word Embeddings for Sentiment Analysis. In *Proceedings of the 2017 Conference on Empirical Methods in Natural Language Processing*, pages 534–539, Copenhagen, Denmark. Association for Computational Linguistics.